

# I/O Branch Bugs on Intermittent Systems

Milijana Surbatovich

Emily Ruppel

## 1 Introduction

## 2 Background

...

### 2.1 A discussion of non-idempotent behavior in intermittent programs

- The program does not have non-idempotent behavior and thus does not have a (non-idempotent) bug.
- The program has non-idempotent behavior. This behavior generally takes two forms: write-after-read dependencies and I/O operations, though there are some less common uses, such as pointer-based stack operations [cite Ratchet]. Bugs off of WAR dependencies have been extensively studied in prior work, and several runtime systems exist to fix them [cite Previous Work]. While there could be other bugs stemming from the use of irrevocable I/O operations, this paper specifically explores memory inconsistencies stemming from branches off of non-idempotent input [Discussed earlier in introduction, presumably].
  - Not all I/O branches potentially lead to buggy behavior – to make memory be in an inconsistent state, non-volatile variables must be written to on at least one side of the branch.
  - Furthermore, on a power fail and re-execution of the non-idempotent region, such a tainted variable has to be the reaching definition for some use, either in expected program execution or along the back-edge introduced by re-executing. A tainted variable can be sanitized by deterministically writing to it from the reset point to the branch, or after the branch before another use [cite figure].
  - Sometimes buggy behavior might only be exposed by re-orderings of instructions during runtime. In such cases where the programmer correctly blocked the tainted variable from

having any uses or wrote the exact same non-volatile write set on both sides of the I/O branch, compiler optimizations might order the initializations on the wrong side of the reset point [This would depend on the runtime system] or change the order of the non-volatile writes such that it is possible to write divergent sets on a power fail. [Note: this needs to be written better.]

In [Figure n], we present a flow chart of the taxonomy of non-idempotent behavior and consequences described above.

## 3 System Overview

The overarching algorithm is fixed-point, iterating through the functions, summarizing new information about the io flow, and iterating through functions again, until there is no new information to be gleaned.

To start with, we don't know where the io enters the program, so all the functions are put in the stack, with no io parameter—the io call/instruction is found internally, based on annotated information in the source code. Once the IO source is found, we call a function that traverses the local data and control flow chain of the io source. In the chain, there are some instructions of interest—namely, branches, which can lead to the bug of mismatched nv write sets, and instructions that result in io tainted values flowing out of the current function that we're examining.

Traversing the data flow primarily uses LLVM's built-in def use chains. From the definition of the tainted variable, we follow the uses in a depth first search, adding on new definitions to the stack. From the traversal function call, we return the list of interesting instructions found. If the instruction was a branch, we examine the non-volatile (i.e.,globals) write sets to see if there is the possibility for the bug. Additionally, the non-volatile variables (used locally and in nested function calls) and local volatile variables written to are control-dependent on the io tainted variable, so we add the globals to the summary information and continue the local traversal with the newly tainted control dependent variables.

For the other types of instructions returned from the traversal, we use the information to add to our list of functions to still be examined to reach fixed point.

Values can leave the local function by:

- being returned from the function
- being passed into a global variable
- being passed as a parameter into another function
- being stored into a call-by-reference parameter (which may be global itself)

If the tainted value was returned from the function *F*, we add all the callers of *F* to the stack, with the value where the return was stored as the tainted io value. (When adding new functions to the examination stack based on summary information, we pass in the tainted value as a parameter instead of searching for the annotations, but then search the chain in the same way. )

Likewise if the tainted value was passed into a global, we find the uses of the global, the parent Function of those uses, and add those to the stack, with the global as the tainted value. If the io was used as a function parameter, we find the callee function and the local version of the argument.

If io was stored into a pass-by-reference parameter, we find all the caller functions and their local value that the argument aliases to. (If the reference para refers to a global, we also add that global to the summary information.)

We add all these new function/tainted value pairs to the stack, and continue to iterate. (We do keep track of what pairs we have already examined, to prevent infinite looping). Eventually, there will be no new information being returned from the traversal function, and we have examined all the functions to which the tainted io has spread.

### 3.1 NV Write-set algorithm

To help with handling arbitrary complexity of control flow within the branches off of io, we precompute the may write sets of each function before hand, using a bottom up examination of the call graph.

What stores we decide to add the write sets does effect what possible bugs will be reported. For instance, one complexity that must be considered is how to handle stores on branches. Say we have a simple program with two nonvolatile variables, *nv0* and *nv1* that branches on an io tainted variable. On one side of the io branch we write to *nv0* and *nv1*. On the other side, we write to *nv0*, but then we branch again. On one side of the branch we write to *nv1*, but on the other side we don't. Thus, in this

example, one execution could have the io branch writing to the same *nv* set on both sides, and another execution doesn't.

There are a few ways to deal with these branches: don't add any stores within the branch basic blocks to the write set, add all the stores to the write set (what we do right now), or make two versions of the write set, adding a different side of the branch to each one. This last option is perhaps most attractive for accuracy, but it is difficult to handle arbitrary complexity, and has a much higher storage overhead.

Another problem with the first two approaches is that if we adjust the above example a little, we can have an io branch where both sides write to *nv0*, and on the second branch, one side also writes to *nv1*. In this case, adding all the stores with the branches detects more bug possibilities, whereas ignoring the branch under reports. So neither of the simpler solutions can be said to definitively under or over report bugs.

### 3.2 Filtering Pass

//TODO

### 3.3 High-level, English explanation of what the tool detects

- The final bug is always the disparate *nv* write sets after a branch off of an io tainted value
- We track the io through explicit (data) and implicit (control) flow
- Tainted io can leave a function through function calls, return values, pass by ref, and globals
- Initial io into the prog has to be annotated (not good or bad, just a thing)

### 3.4 Possible limitations of the tool

- Precomputed write sets
- Overhead?
- How complex can the src code get? – we should prob impose a limitation on branching levels or function depth
- Are there any other suspicious bug patterns? i.e., completeness
- Over or under report bugs? Right now tending towards conservatism

**4 Evaluation**

**5 Implications**

**6 Surprises and Lessons Learned**

**7 Conclusion**